

Machine Learning Exercise Movements

Andrew Washbrun

February 27, 2017

Executive Summary

The Internet of Things allows the classification of previously undocumented activities, such as exercise routines, in gross detail allowing machine learning prediction models to detect which type of exercise a user is attempting. Utilizing data from the Weight Lifting Exercise Dataset, classification and random forest prediction models can predict what type of dumbbell exercise a user is attempting based on accelerometer data gathered on the user's arm, forearm, waist and dumbbell.

Data Wrangling

Accelerometer data is downloaded from dataset's website: <http://groupware.les.inf.puc-rio.br/har>. Training and validation sets are imported and cleaned up to remove uninterpretable values like NA, #DIV/0!, and empty values. Then the first several columns are removed because they contain non-movement related data like user and timestamp. Then the training set is split 70/30 into a training and testing set to prevent model overfitting for the validation set.

```
library(caret); library(rpart);library(rattle)

## Loading required package: lattice
## Loading required package: ggplot2

## Rattle: A free graphical interface for data mining with R.
## Version 4.1.0 Copyright (c) 2006-2015 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.

training <- read.csv('pml-training.csv', na.strings=c("NA", "#DIV/0!", ""))
validation <- read.csv('pml-testing.csv', na.strings=c("NA", "#DIV/0!", ""))

#wrangling data into variables
validation <- validation[, colSums(is.na(training)) == 0]
training <- training[, colSums(is.na(training)) == 0]
training <- training[, -c(1:7)]
validation <- validation[, -c(1:7)]

inTrain <- createDataPartition(y=training$classe,
                               p=0.7, list=FALSE)
trainingData <- training[inTrain,]
testingData <- training[-inTrain,]

names(trainingData)

## [1] "roll_belt"          "pitch_belt"         "yaw_belt"
## [4] "total_accel_belt"   "gyros_belt_x"       "gyros_belt_y"
## [7] "gyros_belt_z"       "accel_belt_x"       "accel_belt_y"
## [10] "accel_belt_z"       "magnet_belt_x"      "magnet_belt_y"
```

```
## [13] "magnet_belt_z"      "roll_arm"          "pitch_arm"
## [16] "yaw_arm"            "total_accel_arm"   "gyros_arm_x"
## [19] "gyros_arm_y"        "gyros_arm_z"       "accel_arm_x"
## [22] "accel_arm_y"        "accel_arm_z"       "magnet_arm_x"
## [25] "magnet_arm_y"       "magnet_arm_z"      "roll_dumbbell"
## [28] "pitch_dumbbell"     "yaw_dumbbell"      "total_accel_dumbbell"
## [31] "gyros_dumbbell_x"   "gyros_dumbbell_y"  "gyros_dumbbell_z"
## [34] "accel_dumbbell_x"   "accel_dumbbell_y"  "accel_dumbbell_z"
## [37] "magnet_dumbbell_x"  "magnet_dumbbell_y" "magnet_dumbbell_z"
## [40] "roll_forearm"       "pitch_forearm"     "yaw_forearm"
## [43] "total_accel_forearm" "gyros_forearm_x"   "gyros_forearm_y"
## [46] "gyros_forearm_z"    "accel_forearm_x"   "accel_forearm_y"
## [49] "accel_forearm_z"    "magnet_forearm_x"  "magnet_forearm_y"
## [52] "magnet_forearm_z"   "classe"
```

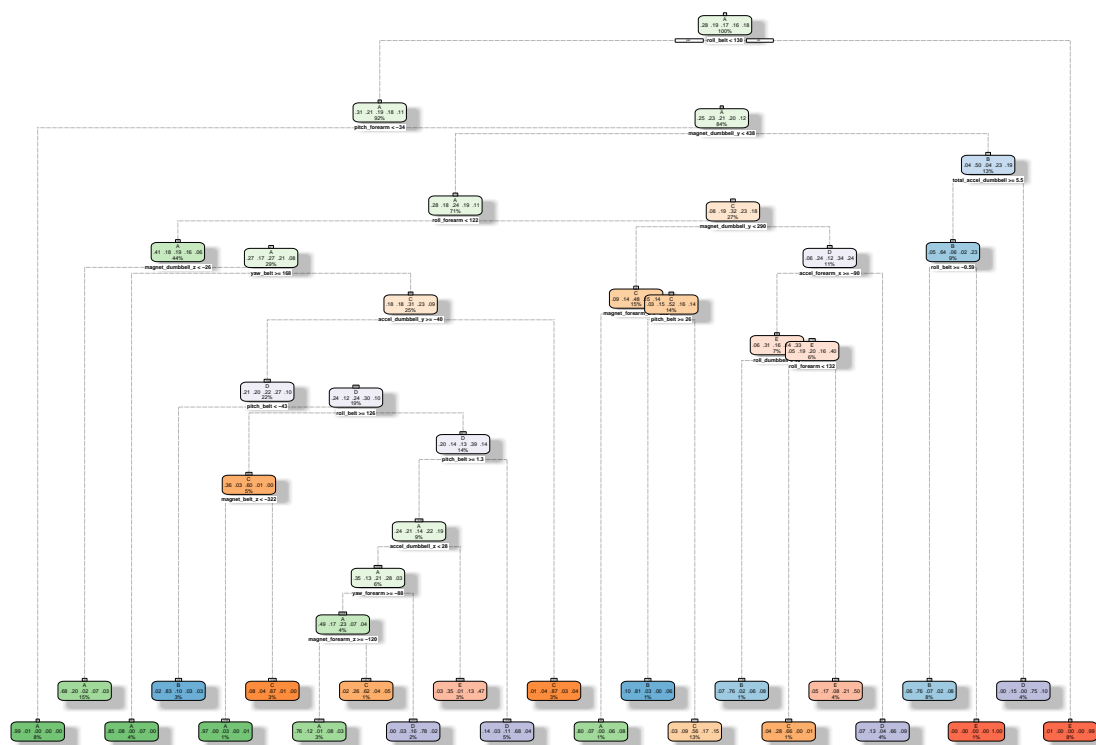
The remaining predictors are accelerometer data from the user and the “classe” variable which is the exercise type.

Recursive Partitioning

Classification trees partition data into logical trees where successive predictors lead to a classification. The classifiers are built such that the root of each branch contains a sufficiently “pure” class, meaning data with identical predictor values will be evaluated to that root class.

```
#tree prediction
modFit <- rpart(classe ~ ., method='class', data=trainingData)
fancyRpartPlot(modFit)
```

```
## Warning: labs do not fit even at cex 0.15, there may be some overplotting
```



Rattle 2017-Feb-27 15:39:16 washbuan

```
predictions <- predict(modFit,testingData,type='class')
confusionStats <- confusionMatrix(predictions,testingData$classe)
confusionStats
```

Confusion Matrix and Statistics

##

Reference

## Prediction	A	B	C	D	E
## A	1512	225	20	98	42
## B	35	582	52	22	57
## C	49	126	858	155	139
## D	47	83	68	621	64
## E	31	123	28	68	780

##

Overall Statistics

##

Accuracy : 0.7397
 ## 95% CI : (0.7283, 0.7509)
 ## No Information Rate : 0.2845
 ## P-Value [Acc > NIR] : < 2.2e-16

##

Kappa : 0.6695
 ## McNemar's Test P-Value : < 2.2e-16

##


Statistics by Class:

##

Class: A Class: B Class: C Class: D Class: E

```
## Sensitivity      0.9032  0.5110  0.8363  0.6442  0.7209
## Specificity     0.9086  0.9650  0.9035  0.9468  0.9479
## Pos Pred Value  0.7970  0.7781  0.6466  0.7033  0.7573
## Neg Pred Value  0.9594  0.8916  0.9631  0.9314  0.9378
## Prevalence      0.2845  0.1935  0.1743  0.1638  0.1839
## Detection Rate  0.2569  0.0989  0.1458  0.1055  0.1325
## Detection Prevalence 0.3223 0.1271 0.2255 0.1500 0.1750
## Balanced Accuracy 0.9059 0.7380 0.8699 0.7955 0.8344
```

```
tableStats <- as.data.frame(confusionStats$table)
ggplot(aes(Prediction,Reference),data=tableStats) +
  geom_tile(aes(fill=Freq)) + scale_fill_gradient(low="green", high="red")
```



The first plot illustrates the classification tree from the top (input) to the bottom (output), from the predictor values to the exercise outcome. The second output is a set of statistics for the model. The main takeaway is the accuracy: 71.5%. Not especially accurate so we'll employ another model. The last plot is an illustration of the Confusion Matrix statistics.

Random Forest

Random forests extend the idea of classification trees with random bootstrapping. By resampling and averaging models, a more robust decision tree is created.

```
#Random Forest
library(randomForest)
```

```
## randomForest 4.6-12
```

```

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'


## The following object is masked from 'package:ggplot2':
##
##     margin

forestModel <- randomForest(classe~.,data=trainingData)
forestPredictions <- predict(forestModel,testingData,type='class')
forestConfStats <- confusionMatrix(forestPredictions,testingData$classe)
forestConfStats

## Confusion Matrix and Statistics
##
##           Reference
## Prediction    A    B    C    D    E
##           A 1670    8    0    0    0
##           B    3 1129    3    0    0
##           C    0    2 1022    5    0
##           D    0    0    1  959    2
##           E    1    0    0    0 1080
##
## Overall Statistics
##
##           Accuracy : 0.9958
##           95% CI : (0.9937, 0.9972)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9946
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9976  0.9912  0.9961  0.9948  0.9982
## Specificity      0.9981  0.9987  0.9986  0.9994  0.9998
## Pos Pred Value   0.9952  0.9947  0.9932  0.9969  0.9991
## Neg Pred Value   0.9990  0.9979  0.9992  0.9990  0.9996
## Prevalence       0.2845  0.1935  0.1743  0.1638  0.1839
## Detection Rate   0.2838  0.1918  0.1737  0.1630  0.1835
## Detection Prevalence 0.2851  0.1929  0.1749  0.1635  0.1837
## Balanced Accuracy 0.9979  0.9950  0.9973  0.9971  0.9990

forestTableStats <- as.data.frame(forestConfStats$table)
ggplot(aes(Prediction,Reference),data=forestTableStats) +
  geom_tile(aes(fill=Freq)) + scale_fill_gradient(low="green", high="red")

```



Comparing the accuracy between random forest and the standard classification tree model: 99.3% vs. 71.5% the random forest model is a more accurate model for prediction.